

## Practical 4 PART A and B

B-76

Shivam Gupta

### ▼ PART(A): Linear Regression

```
# 1. Consider a set of points x and the predicted values y as given below:
```

```
# x y
# 0 2
# 2 4
# 3 5
# 5 6
# 6 7.5
# 7 8.5
# 8 9
# 8 10
# 9 11
# 10 12
```

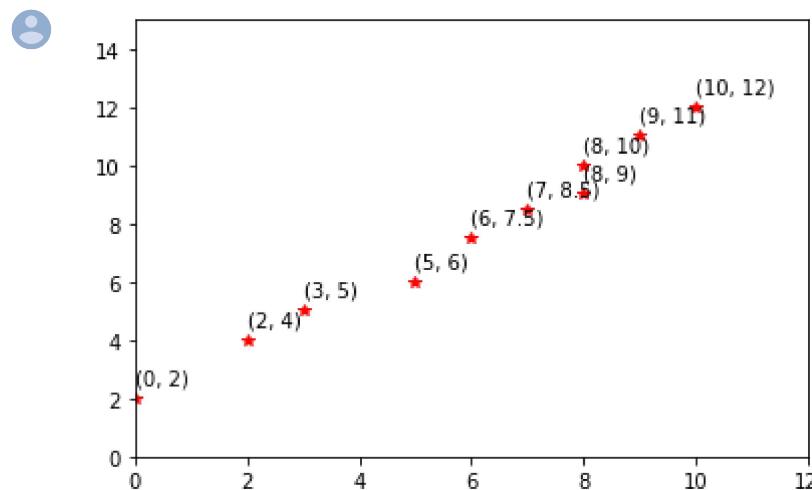
```
# (a) Plot the data.
```

```
x_org=[0,2,3,5,6,7,8,8,9,10]
y_org=[2,4,5,6,7.5,8.5,9,10,11,12]
```

```
import matplotlib.pyplot as plt
plt.plot(x_org, y_org, 'r*')
plt.axis([0, 12, 0, 15])

for i, j in zip(x_org, y_org):
    plt.text(i, j+0.5, '({}, {})'.format(i, j))

plt.show()
```



```
# (b) Use LinearRegression model from sklearn library to perform linear regression.

import numpy as np
from sklearn.linear_model import LinearRegression
x = np.array(x_org).reshape((-1, 1))
y = np.array(y_org)
print('x is',x)
print('y is',y)

model = LinearRegression()
model.fit(x, y)

r_sq = model.score(x, y)
print(f"coefficient of determination: {r_sq}")

print(f"intercept: {model.intercept_}")

print(f"slope: {model.coef_}")

x is [[ 0]
 [ 2]
 [ 3]
 [ 5]
 [ 6]
 [ 7]
 [ 8]
 [ 8]
 [ 9]
 [10]]
y is [ 2.   4.   5.   6.   7.5  8.5  9.   10.  11.  12. ]
coefficient of determination: 0.9832917172807847
intercept: 1.8274058577405858
slope: [0.97803347]
```

# (c) Consider a data value for x and predict the value of y using the above model.

```
y_pred = model.predict(x)
print(f"predicted response:\n{y_pred}")

predicted response:
[ 1.82740586  3.7834728   4.76150628  6.71757322  7.69560669  8.67364017
 9.65167364  9.65167364 10.62970711 11.60774059]
```

# 2. The given table shows the midterm and final exam grades obtained for students in a database course.

# X [midterm exam] Y [Final Exam]

# 72 84

# 50 63

# 81 77

# 74 78

# 94 90

# 86 75

# 59 49

# 83 79

```
# 65 77
# 33 52
# 88 74
# 81 90

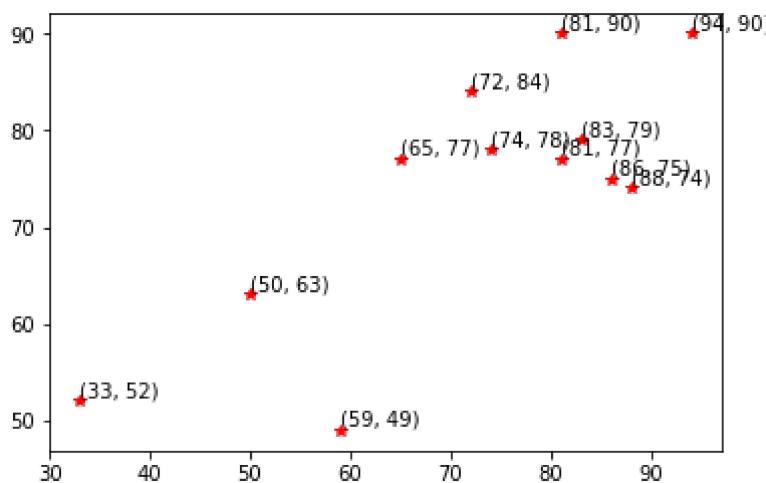
# (a) Plot the data. Do x and y seem to have a linear relationship?

x=[72,50,81,74,94,86,59,83,65,33,88,81]
y=[84,63,77,78,90,75,49,79,77,52,74,90]

plt.plot(x, y, 'r*')

for i, j in zip(x, y):
    plt.text(i, j+0.5, '({}, {})'.format(i, j))

plt.show()
```



Not linear

```
# (b) Use the method of least squares to find an equation for the prediction of a student's
# final exam grade based on the student's midterm grade in the course.
# Write a function in python to compute the coefficients and equation. [Do not use the
# inbuild library method.]
```

```
n=len(x)
sumx=sum(x)
sumy=sum(y)
sumxy=0
xsqsum=0

for i in range(n):
    sumxy+=x[i]*y[i]
    xsqsum+=x[i]*x[i]
sumxsq=sumx*sumx

m=(n*sumxy-(sumx)*(sumy))/(n*xsqsum-sumxsq)

print("Slope: m = ",m)

xmean=np.mean(x)
```

```
ymean=np.mean(y)
```

```
c=ymean-xmean*m
```

```
print("Y-intercept: c=",c)
```

```
Slope: m = 0.5816000773918932
```

```
Y-intercept: c= 32.027861081551706
```

```
# (c) Also show the plot with the datapoints and the obtained linear equation line.
```

```
x=[72,50,81,74,94,86,59,83,65,33,88,81]
```

```
y=[84,63,77,78,90,75,49,79,77,52,74,90]
```

```
plt.plot(x, y, 'r*')
```

```
for i, j in zip(x, y):
```

```
    plt.text(i, j+0.5, '({}, {})'.format(i, j))
```

```
x = np.array(x)
```

```
y = m*x+c
```

```
plt.plot(x, y, '-r', label='y=0.58 x + 31.99')
```

```
plt.title(f'Graph of y=0.58 x + 31.99')
```

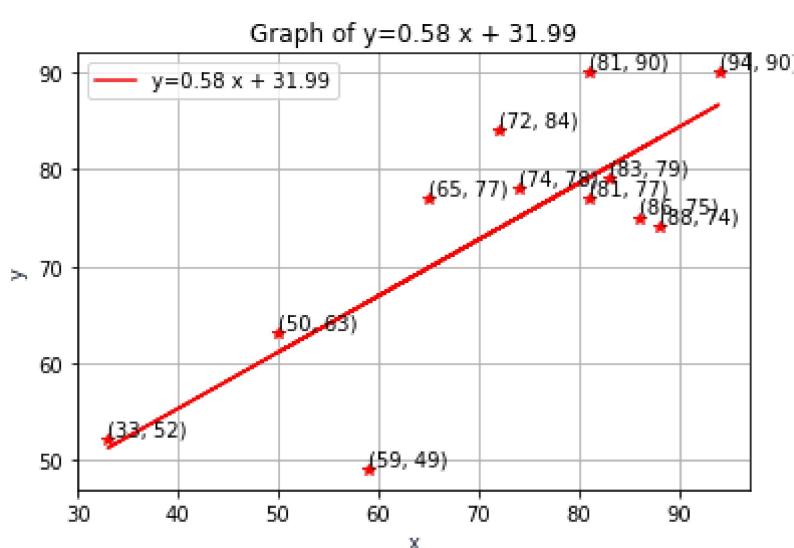
```
plt.xlabel('x', color='#1C2833')
```

```
plt.ylabel('y', color='#1C2833')
```

```
plt.legend(loc='upper left')
```

```
plt.grid()
```

```
plt.show()
```



```
# (d) Predict the final exam grade of student who received an 86 in the midterm exam  
# based on the equation of least squares.
```

```
print("final exam grade of student who received an 86 in the midterm exam is ",m*86+c)
```

```
final exam grade of student who received an 86 in the midterm exam is 82.0454677372
```

```
# Do as Directed [Multiple Linear regression]:
```

```
# 3. Perform Multiple Linear regression on cars.csv dataset.  
# (a) Analyse each column of the datset using appropriate visualization technique.
```

```
import pandas as pd  
import matplotlib.pyplot as mp  
import seaborn as sns # for visualization  
import matplotlib.pyplot as plt  
  
# take data  
df = pd.read_csv("cars.csv")  
print(df)  
  
c = df.corr()  
print("\n\n",c)  
sns.heatmap(c,annot=True,cmap='Blues')  
plt.show()
```

|    | Car        | Model      | Volume | Weight | CO2 |
|----|------------|------------|--------|--------|-----|
| 0  | Toyoto     | Aygo       | 1000   | 790    | 99  |
| 1  | Mitsubishi | Space Star | 1200   | 1160   | 95  |
| 2  | Skoda      | Citigo     | 1000   | 929    | 95  |
| 3  | Fiat       | 500        | 900    | 865    | 90  |
| 4  | Mini       | Cooper     | 1500   | 1140   | 105 |
| 5  | VW         | Up!        | 1000   | 929    | 105 |
| 6  | Skoda      | Fabia      | 1400   | 1109   | 90  |
| 7  | Mercedes   | A-Class    | 1500   | 1365   | 92  |
| 8  | Ford       | Fiesta     | 1500   | 1112   | 98  |
| 9  | Audi       | A1         | 1600   | 1150   | 99  |
| 10 | Hyundai    | I20        | 1100   | 980    | 99  |
| 11 | Suzuki     | Swift      | 1300   | 990    | 101 |
| 12 | Ford       | Fiesta     | 1000   | 1112   | 99  |
| 13 | Honda      | Civic      | 1600   | 1252   | 94  |
| 14 | Hyundai    | I30        | 1600   | 1326   | 97  |

```
# (b) Consider only weight and volume columns as multiple variables to predict the
# CO2 emission.
```

```
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn import linear_model

shape=df.shape
print(shape,"\\n")
print(df.columns,"\\n")
print(df.describe(),"\\n\\n")

# splitting Dataset in Dependent & Independent Variables
X = df[['Weight','Volume']]
y = df['CO2']

print("Independent Variables: \\n\\n",X)
print("Dependent Variables: \\n\\n",y)
```

```
x_train,x_test,y_train,y_test = train_test_split(X,y,train_size=0.7,random_state=0)
```

```
regr = linear_model.LinearRegression()
regr.fit(X, y)
print('\\nModel has been trained successfully')
```

```
y_pred = regr.predict(x_test)
print("\\nPredicted CO2:\\n",y_pred)
```

```
(36, 5)
```

```
Index(['Car', 'Model', 'Volume', 'Weight', 'CO2'], dtype='object')
```

|       | Volume      | Weight      | CO2        |
|-------|-------------|-------------|------------|
| count | 36.000000   | 36.000000   | 36.000000  |
| mean  | 1611.111111 | 1292.277778 | 102.027778 |
| std   | 388.975047  | 242.123889  | 7.454571   |
| min   | 900.000000  | 790.000000  | 90.000000  |
| 25%   | 1475.000000 | 1117.250000 | 97.750000  |
| 50%   | 1600.000000 | 1329.000000 | 99.000000  |

|     |             |             |            |
|-----|-------------|-------------|------------|
| 75% | 2000.000000 | 1418.250000 | 105.000000 |
| max | 2500.000000 | 1746.000000 | 120.000000 |

### Independent Variables:

|    | Weight | Volume |
|----|--------|--------|
| 0  | 790    | 1000   |
| 1  | 1160   | 1200   |
| 2  | 929    | 1000   |
| 3  | 865    | 900    |
| 4  | 1140   | 1500   |
| 5  | 929    | 1000   |
| 6  | 1109   | 1400   |
| 7  | 1365   | 1500   |
| 8  | 1112   | 1500   |
| 9  | 1150   | 1600   |
| 10 | 980    | 1100   |
| 11 | 990    | 1300   |
| 12 | 1112   | 1000   |
| 13 | 1252   | 1600   |
| 14 | 1326   | 1600   |
| 15 | 1330   | 1600   |
| 16 | 1365   | 1600   |
| 17 | 1280   | 2200   |
| 18 | 1119   | 1600   |
| 19 | 1328   | 2000   |
| 20 | 1584   | 1600   |
| 21 | 1428   | 2000   |
| 22 | 1365   | 2100   |
| 23 | 1415   | 1600   |
| 24 | 1415   | 2000   |
| 25 | 1465   | 1500   |
| 26 | 1490   | 2000   |
| 27 | 1725   | 2000   |
| 28 | 1523   | 1600   |
| 29 | 1705   | 2000   |
| 30 | 1605   | 2100   |
| 31 | 1746   | 2000   |
| 32 | 1235   | 1600   |
| 33 | 1390   | 1600   |
| 34 | 1405   | 1600   |
| 35 | 1395   | 2500   |

### Dependent Variables:

|   |    |
|---|----|
| 0 | 99 |
| - | -- |

```
# (c) Predict the CO2 emission of a car where the weight is 2300kg, and the volume is
# 1300ccm. Also conclude about the obtained result.
```

```
x_train,x_test,y_train,y_test = train_test_split(X,y,train_size=0.7,random_state=0)

reg = linear_model.LinearRegression()
reg.fit(X, y)

prediction = reg.predict([[2300, 1300]])
print("\n\nCO2 emission of car for weight=2300kg, volume=1300cm3 is => ",prediction,"\\n\\n")
```

```
CO2 emission of car for weight=2300kg, volume=1300cm3 is => [107.2087328]
```

```
C:\Users\DELL\AppData\Roaming\Python\Python39\site-packages\sklearn\base.py:450: User
warnings.warn(
```

```
# (d) Perform prediction on few more values.
```

```
prediction = reg.predict([[1300, 300]])
print("\nCO2 emission of car for weight=2300kg, volume=1300cm3 is => ", prediction, "\n")
```

```
prediction = reg.predict([[200, 130]])
print("\nCO2 emission of car for weight=2300kg, volume=1300cm3 is => ", prediction, "\n")
```

```
prediction = reg.predict([[330, 83]])
print("\nCO2 emission of car for weight=2300kg, volume=1300cm3 is => ", prediction, "\n")
```

```
CO2 emission of car for weight=2300kg, volume=1300cm3 is => [91.852528]
```

```
CO2 emission of car for weight=2300kg, volume=1300cm3 is => [82.21959222]
```

```
CO2 emission of car for weight=2300kg, volume=1300cm3 is => [82.83436827]
```

```
C:\Users\DELL\AppData\Roaming\Python\Python39\site-packages\sklearn\base.py:450: User
warnings.warn(
```

```
C:\Users\DELL\AppData\Roaming\Python\Python39\site-packages\sklearn\base.py:450: User
warnings.warn(
```

```
C:\Users\DELL\AppData\Roaming\Python\Python39\site-packages\sklearn\base.py:450: User
warnings.warn(
```

```
# (e) Show the coefficient obtained and conclude.
```

```
print("Coefficient obtained of multiple regression are = ", reg.coef_)
```

```
Coefficient obtained of multiple regression are = [0.00755095 0.00780526]
```

```
# 4. Perform linear regression on the dataset [use dataset: kc_house_data.csv]
```

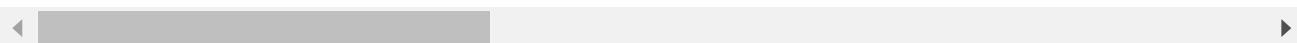
```
# (a) Load the dataset, display it, visualize various columns and explain the dataset
# composition.
```

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
df = pd.read_csv("kc_house_data.csv")
df
```

|              | <b>id</b>  | <b>date</b>     | <b>price</b> | <b>bedrooms</b> | <b>bathrooms</b> | <b>sqft_living</b> | <b>sqft</b> |
|--------------|------------|-----------------|--------------|-----------------|------------------|--------------------|-------------|
| <b>0</b>     | 7129300520 | 20141013T000000 | 221900.0     | 3               | 1.00             | 1180               | 1           |
| <b>1</b>     | 6414100192 | 20141209T000000 | 538000.0     | 3               | 2.25             | 2570               | 1           |
| <b>2</b>     | 5631500400 | 20150225T000000 | 180000.0     | 2               | 1.00             | 770                | 10          |
| <b>3</b>     | 2487200875 | 20141209T000000 | 604000.0     | 4               | 3.00             | 1960               | 1           |
| <b>4</b>     | 1954400510 | 20150218T000000 | 510000.0     | 3               | 2.00             | 1680               | 1           |
| ...          | ...        | ...             | ...          | ...             | ...              | ...                | ...         |
| <b>21608</b> | 263000018  | 20140521T000000 | 360000.0     | 3               | 2.50             | 1530               | 1           |
| <b>21609</b> | 6600060120 | 20150223T000000 | 400000.0     | 4               | 2.50             | 2310               | 1           |
| <b>21610</b> | 1523300141 | 20140623T000000 | 402101.0     | 2               | 0.75             | 1020               | 1           |
| <b>21611</b> | 291310100  | 20150116T000000 | 400000.0     | 3               | 2.50             | 1600               | 1           |
| <b>21612</b> | 1523300157 | 20141015T000000 | 325000.0     | 2               | 0.75             | 1020               | 1           |

21613 rows × 21 columns



```
print("Summary of the DataFrame:\n")
df.info()
```

Summary of the DataFrame:

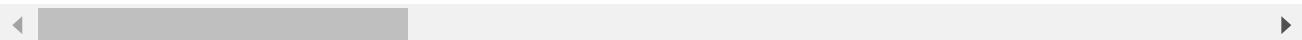
```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21613 entries, 0 to 21612
Data columns (total 21 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   id               21613 non-null   int64  
 1   date              21613 non-null   object 
 2   price             21613 non-null   float64 
 3   bedrooms          21613 non-null   int64  
 4   bathrooms         21613 non-null   float64 
 5   sqft_living       21613 non-null   int64  
 6   sqft_lot          21613 non-null   int64  
 7   floors            21613 non-null   float64 
 8   waterfront        21613 non-null   int64  
 9   view               21613 non-null   int64  
 10  condition         21613 non-null   int64  
 11  grade              21613 non-null   int64  
 12  sqft_above        21613 non-null   int64  
 13  sqft_basement     21613 non-null   int64  
 14  yr_built          21613 non-null   int64  
 15  yr_renovated      21613 non-null   int64  
 16  zipcode            21613 non-null   int64  
 17  lat                21613 non-null   float64 
 18  long               21613 non-null   float64 
 19  sqft_living15     21613 non-null   int64  
 20  sqft_lot15         21613 non-null   int64 
```

```
dtypes: float64(5), int64(15), object(1)
memory usage: 3.5+ MB
```

```
# (b) Describe the dataset
```

```
df.describe()
```

|              | <b>id</b>    | <b>price</b> | <b>bedrooms</b> | <b>bathrooms</b> | <b>sqft_living</b> | <b>sqft_living15</b> |
|--------------|--------------|--------------|-----------------|------------------|--------------------|----------------------|
| <b>count</b> | 2.161300e+04 | 2.161300e+04 | 21613.000000    | 21613.000000     | 21613.000000       | 2.161300e+04         |
| <b>mean</b>  | 4.580302e+09 | 5.400881e+05 | 3.370842        | 2.114757         | 2079.899736        | 1.510697e+09         |
| <b>std</b>   | 2.876566e+09 | 3.671272e+05 | 0.930062        | 0.770163         | 918.440897         | 4.142051e+08         |
| <b>min</b>   | 1.000102e+06 | 7.500000e+04 | 0.000000        | 0.000000         | 290.000000         | 5.200000e+04         |
| <b>25%</b>   | 2.123049e+09 | 3.219500e+05 | 3.000000        | 1.750000         | 1427.000000        | 5.040000e+08         |
| <b>50%</b>   | 3.904930e+09 | 4.500000e+05 | 3.000000        | 2.250000         | 1910.000000        | 7.618000e+08         |
| <b>75%</b>   | 7.308900e+09 | 6.450000e+05 | 4.000000        | 2.500000         | 2550.000000        | 1.068800e+09         |
| <b>max</b>   | 9.900000e+09 | 7.700000e+06 | 33.000000       | 8.000000         | 13540.000000       | 1.651359e+09         |



```
# (c) Houses with how many bedrooms are most sold?
```

```
print("Houses with 3 Bedrooms are most sold.")
df['bedrooms'].value_counts()
```

```
Houses with 3 Bedrooms are most sold.
```

|    |      |
|----|------|
| 3  | 9824 |
| 4  | 6882 |
| 2  | 2760 |
| 5  | 1601 |
| 6  | 272  |
| 1  | 199  |
| 7  | 38   |
| 0  | 13   |
| 8  | 13   |
| 9  | 6    |
| 10 | 3    |
| 11 | 1    |
| 33 | 1    |

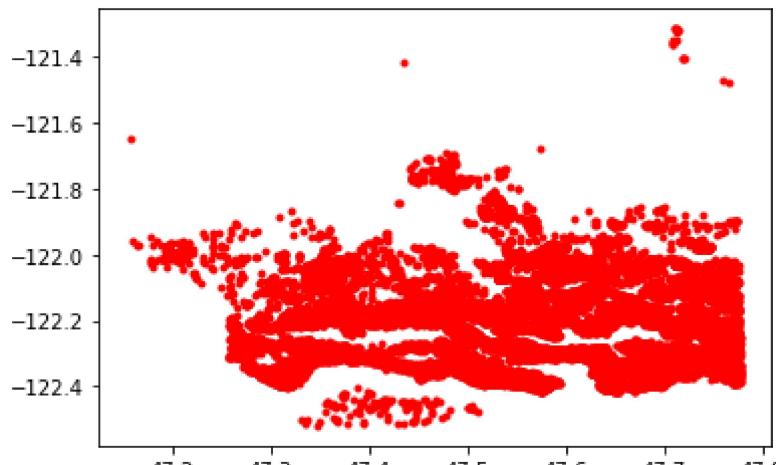
```
Name: bedrooms, dtype: int64
```

```
# (d) Visualizing the location of the houses based on latitude and longitude. Use
# appropriate plot.
```

```
lati=df['lat']
longi=df['long']
```



```
plt.plot(lati, longi, 'r.')
plt.show()
```



```
# (e) Find the correlated features/ columns
```

```
c = df.corr()
print("\n",c)
sns.heatmap(c,cmap='Blues')
plt.show()
```

|                      | <b>id</b> | <b>price</b> | <b>bedrooms</b> | <b>bathrooms</b> | <b>sqft_living</b> | <b>sqft_lot</b> | \ |
|----------------------|-----------|--------------|-----------------|------------------|--------------------|-----------------|---|
| <b>id</b>            | 1.000000  | -0.016762    | 0.001286        | 0.005160         | -0.012258          | -0.132109       |   |
| <b>price</b>         | -0.016762 | 1.000000     | 0.308350        | 0.525138         | 0.702035           | 0.089661        |   |
| <b>bedrooms</b>      | 0.001286  | 0.308350     | 1.000000        | 0.515884         | 0.576671           | 0.031703        |   |
| <b>bathrooms</b>     | 0.005160  | 0.525138     | 0.515884        | 1.000000         | 0.754665           | 0.087740        |   |
| <b>sqft_living</b>   | -0.012258 | 0.702035     | 0.576671        | 0.754665         | 1.000000           | 0.172826        |   |
| <b>sqft_lot</b>      | -0.132109 | 0.089661     | 0.031703        | 0.087740         | 0.172826           | 1.000000        |   |
| <b>floors</b>        | 0.018525  | 0.256794     | 0.175429        | 0.500653         | 0.353949           | -0.005201       |   |
| <b>waterfront</b>    | -0.002721 | 0.266369     | -0.006582       | 0.063744         | 0.103818           | 0.021604        |   |
| <b>view</b>          | 0.011592  | 0.397293     | 0.079532        | 0.187737         | 0.284611           | 0.074710        |   |
| <b>condition</b>     | -0.023783 | 0.036362     | 0.028472        | -0.124982        | -0.058753          | -0.008958       |   |
| <b>grade</b>         | 0.008130  | 0.667434     | 0.356967        | 0.664983         | 0.762704           | 0.113621        |   |
| <b>sqft_above</b>    | -0.010842 | 0.605567     | 0.477600        | 0.685342         | 0.876597           | 0.183512        |   |
| <b>sqft_basement</b> | -0.005151 | 0.323816     | 0.303093        | 0.283770         | 0.435043           | 0.015286        |   |
| <b>yr_built</b>      | 0.021380  | 0.054012     | 0.154178        | 0.506019         | 0.318049           | 0.053080        |   |
| <b>yr_renovated</b>  | -0.016907 | 0.126434     | 0.018841        | 0.050739         | 0.055363           | 0.007644        |   |
| <b>zipcode</b>       | -0.008224 | -0.053203    | -0.152668       | -0.203866        | -0.199430          | -0.129574       |   |
| <b>lat</b>           | -0.001891 | 0.307003     | -0.008931       | 0.024573         | 0.052529           | -0.085683       |   |
| <b>long</b>          | 0.020799  | 0.021626     | 0.129473        | 0.223042         | 0.240223           | 0.229521        |   |
| <b>sqft_living15</b> | -0.002901 | 0.585379     | 0.391638        | 0.568634         | 0.756420           | 0.144608        |   |
| <b>sqft_lot15</b>    | -0.138798 | 0.082447     | 0.029244        | 0.087175         | 0.183286           | 0.718557        |   |

```
# (f) Find null values and fill with mean value for all columns
```

```
l=[]
for i in df.columns:
    l.append(i)
l=l[2:]

print(l)

for col in l:
    mean=df[col].mean()
    df[col].fillna(value=mean,inplace=True)

df
```

|       | id         | date            | price    | bedrooms | bathrooms | sqft_living | sqft_lot | floors | waterfront | view | condition | grade | sqft_above | sqft_basement | yr_built | yr_renovated | zipcode | lat     | long     |          |
|-------|------------|-----------------|----------|----------|-----------|-------------|----------|--------|------------|------|-----------|-------|------------|---------------|----------|--------------|---------|---------|----------|----------|
| 0     | 7129300520 | 20141013T000000 | 221900.0 |          | 3         | 1.00        | 1180     | 5650   | 1.0        | 0    | 0         | 3     | 7          | 1180          | 1955     | 0            | 98178   | 47.5112 | -122.257 |          |
| 1     | 6414100192 | 20141209T000000 | 538000.0 |          | 3         | 2.25        | 2570     | 7242   | 2.0        | 0    | 0         | 3     | 7          | 2170          | 1951     | 1991         | 98125   | 47.7210 | -122.319 |          |
| 2     | 5631500400 | 20150225T000000 | 180000.0 |          | 2         | 1.00        | 770      | 10000  | 1.0        | 0    | 0         | 4     | 4          | 3.00          | 1960     | 5000         | 0       | 98028   | 47.7379  | -122.233 |
| 3     |            |                 |          |          |           |             |          |        |            |      |           | 3     | 3          | 2.00          | 1680     | 8080         | 1.0     | 98136   | 47.5208  | -122.393 |
| 4     |            |                 |          |          |           |             |          |        |            |      |           | 3     | 3          | 2.00          | 1680     | 8080         | 1.0     | 98074   | 47.6168  | -122.045 |
| ...   |            |                 |          |          |           |             |          |        |            |      |           | ...   | ...        | ...           | ...      | ...          | ...     | ...     | ...      |          |
| 21608 |            |                 |          |          |           |             |          |        |            |      |           | 3     | 3          | 2.50          | 1530     | 1131         | 3.0     | 98103   | 47.6993  | -122.346 |
| 21609 |            |                 |          |          |           |             |          |        |            |      |           | 4     | 4          | 2.50          | 2310     | 5813         | 2.0     | 98146   | 47.5107  | -122.362 |
| 21610 |            |                 |          |          |           |             |          |        |            |      |           | 2     | 2          | 0.75          | 1020     | 1350         | 2.0     | 98144   | 47.5944  | -122.299 |
| 21611 |            |                 |          |          |           |             |          |        |            |      |           | 3     | 3          | 2.50          | 1600     | 2388         | 2.0     | 98027   | 47.5345  | -122.069 |
| 21612 |            |                 |          |          |           |             |          |        |            |      |           | 2     | 2          | 0.75          | 1020     | 1076         | 2.0     | 98144   | 47.5941  | -122.299 |

```
# (g) Find dependant and independent data (place in X and y)
```

```
X = df[1[1:]]
y = df['price']
```

```
print("Independent Variables: \n\n",X)
print("Dependent Variables: \n\n",y)
```

Independent Variables:

|       | bedrooms | bathrooms | sqft_living | sqft_lot | floors | waterfront | view | \   |
|-------|----------|-----------|-------------|----------|--------|------------|------|-----|
| 0     | 3        | 1.00      | 1180        | 5650     | 1.0    | 0          | 0    |     |
| 1     | 3        | 2.25      | 2570        | 7242     | 2.0    | 0          | 0    |     |
| 2     | 2        | 1.00      | 770         | 10000    | 1.0    | 0          | 0    |     |
| 3     | 4        | 3.00      | 1960        | 5000     | 1.0    | 0          | 0    |     |
| 4     | 3        | 2.00      | 1680        | 8080     | 1.0    | 0          | 0    |     |
| ...   | ...      | ...       | ...         | ...      | ...    | ...        | ...  | ... |
| 21608 | 3        | 2.50      | 1530        | 1131     | 3.0    | 0          | 0    |     |
| 21609 | 4        | 2.50      | 2310        | 5813     | 2.0    | 0          | 0    |     |
| 21610 | 2        | 0.75      | 1020        | 1350     | 2.0    | 0          | 0    |     |
| 21611 | 3        | 2.50      | 1600        | 2388     | 2.0    | 0          | 0    |     |
| 21612 | 2        | 0.75      | 1020        | 1076     | 2.0    | 0          | 0    |     |

|       | condition | grade | sqft_above | sqft_basement | yr_built | yr_renovated | \   |
|-------|-----------|-------|------------|---------------|----------|--------------|-----|
| 0     | 3         | 7     | 1180       | 0             | 1955     | 0            |     |
| 1     | 3         | 7     | 2170       | 400           | 1951     | 1991         |     |
| 2     | 3         | 6     | 770        | 0             | 1933     | 0            |     |
| 3     | 5         | 7     | 1050       | 910           | 1965     | 0            |     |
| 4     | 3         | 8     | 1680       | 0             | 1987     | 0            |     |
| ...   | ...       | ...   | ...        | ...           | ...      | ...          | ... |
| 21608 | 3         | 8     | 1530       | 0             | 2009     | 0            |     |
| 21609 | 3         | 8     | 2310       | 0             | 2014     | 0            |     |
| 21610 | 3         | 7     | 1020       | 0             | 2009     | 0            |     |
| 21611 | 3         | 8     | 1600       | 0             | 2004     | 0            |     |
| 21612 | 3         | 7     | 1020       | 0             | 2008     | 0            |     |

|       | zipcode | lat     | long     | sqft_living15 | sqft_lot15 |
|-------|---------|---------|----------|---------------|------------|
| 0     | 98178   | 47.5112 | -122.257 | 1340          | 5650       |
| 1     | 98125   | 47.7210 | -122.319 | 1690          | 7639       |
| 2     | 98028   | 47.7379 | -122.233 | 2720          | 8062       |
| 3     | 98136   | 47.5208 | -122.393 | 1360          | 5000       |
| 4     | 98074   | 47.6168 | -122.045 | 1800          | 7503       |
| ...   | ...     | ...     | ...      | ...           | ...        |
| 21608 | 98103   | 47.6993 | -122.346 | 1530          | 1509       |
| 21609 | 98146   | 47.5107 | -122.362 | 1830          | 7200       |
| 21610 | 98144   | 47.5944 | -122.299 | 1020          | 2007       |
| 21611 | 98027   | 47.5345 | -122.069 | 1410          | 1287       |
| 21612 | 98144   | 47.5941 | -122.299 | 1020          | 1357       |

[21613 rows x 18 columns]

## Dependent Variables:

```

0      221900.0
1      538000.0
2      180000.0
3      604000.0
4      510000.0
...
21608   360000.0
21609   400000.0
21610   402101.0
21611   400000.0
21612   325000.0
Name: price, Length: 21613, dtype: float64

```

## # (h) Split train and test data

```
x_train,x_test,y_train,y_test = train_test_split(X,y,train_size=0.7,random_state=0)
print(x_test)
print(y_test)
```

|       | bedrooms | bathrooms | sqft_living | sqft_lot | floors | waterfront | view | \ |
|-------|----------|-----------|-------------|----------|--------|------------|------|---|
| 17384 | 2        | 1.50      | 1430        | 1650     | 3.0    | 0          | 0    |   |
| 722   | 4        | 3.25      | 4670        | 51836    | 2.0    | 0          | 0    |   |
| 2680  | 2        | 0.75      | 1440        | 3700     | 1.0    | 0          | 0    |   |
| 18754 | 2        | 1.00      | 1130        | 2640     | 1.0    | 0          | 0    |   |
| 14554 | 4        | 2.50      | 3180        | 9603     | 2.0    | 0          | 2    |   |
| ...   | ...      | ...       | ...         | ...      | ...    | ...        | ...  |   |
| 18588 | 3        | 2.25      | 1560        | 8570     | 1.0    | 0          | 0    |   |
| 6784  | 4        | 1.75      | 2360        | 6000     | 1.0    | 0          | 0    |   |
| 14510 | 3        | 2.50      | 2150        | 25705    | 1.5    | 0          | 0    |   |
| 18917 | 3        | 1.75      | 1480        | 8009     | 1.0    | 0          | 0    |   |
| 11750 | 3        | 1.00      | 1320        | 9239     | 1.0    | 0          | 0    |   |

|       | condition | grade | sqft_above | sqft_basement | yr_built | yr_renovated | \ |
|-------|-----------|-------|------------|---------------|----------|--------------|---|
| 17384 | 3         | 7     | 1430       | 0             | 1999     | 0            |   |
| 722   | 4         | 12    | 4670       | 0             | 1988     | 0            |   |
| 2680  | 3         | 7     | 1200       | 240           | 1914     | 0            |   |
| 18754 | 4         | 8     | 1130       | 0             | 1927     | 0            |   |
| 14554 | 3         | 9     | 3180       | 0             | 2002     | 0            |   |
| ...   | ...       | ...   | ...        | ...           | ...      | ...          |   |
| 18588 | 5         | 7     | 1080       | 480           | 1977     | 0            |   |
| 6784  | 3         | 7     | 1280       | 1080          | 1955     | 0            |   |
| 14510 | 3         | 6     | 2150       | 0             | 1980     | 2009         |   |
| 18917 | 3         | 7     | 980        | 500           | 1980     | 0            |   |
| 11750 | 4         | 7     | 1320       | 0             | 1968     | 0            |   |

|       | zipcode | lat     | long     | sqft_living15 | sqft_lot15 |
|-------|---------|---------|----------|---------------|------------|
| 17384 | 98125   | 47.7222 | -122.290 | 1430          | 1650       |
| 722   | 98005   | 47.6350 | -122.164 | 4230          | 41075      |
| 2680  | 98107   | 47.6707 | -122.364 | 1440          | 4300       |
| 18754 | 98109   | 47.6438 | -122.357 | 1680          | 3200       |
| 14554 | 98155   | 47.7717 | -122.277 | 2440          | 15261      |
| ...   | ...     | ...     | ...      | ...           | ...        |
| 18588 | 98004   | 47.6155 | -122.210 | 2660          | 9621       |
| 6784  | 98109   | 47.6465 | -122.357 | 1700          | 3460       |
| 14510 | 98058   | 47.4514 | -122.089 | 1850          | 20160      |
| 18917 | 98032   | 47.3657 | -122.280 | 1790          | 7678       |

```
11750    98092  47.3120 -122.183          1320        8415
[6484 rows x 18 columns]
17384    297000.0
722      1578000.0
2680     562100.0
18754    631500.0
14554    780000.0
...
18588    1100000.0
6784     700000.0
14510    380000.0
18917    268000.0
11750    206000.0
Name: price, Length: 6484, dtype: float64
```

```
# (i) Train the model and test it. Find the accuracy.
```

```
reg = linear_model.LinearRegression()
reg.fit(X, y)

print('\nModel has been trained successfully')
print('R-squared (testing) ', reg.score(x_test, y_test))
```

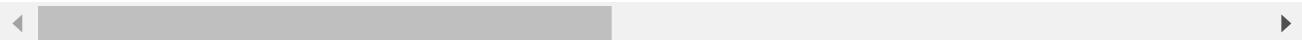
```
Model has been trained successfully
R-squared (testing)  0.689837634636076
```

```
# (j) Test the model using some arbitrary input.
```

```
# 'bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors', 'waterfront', 'view', 'con
prediction = reg.predict([[4,1.5,1430,1650,3,0,0,3,7,1430,0,2000,0,98125, 47.7,-122.5,1330
print("\nHouse Price Prediction for arbitrary input is => ",prediction,"\\n")
```

```
House Price Prediction for arbitrary input is => [332067.18606665]
```

```
C:\Users\DELL\AppData\Roaming\Python\Python39\site-packages\sklearn\base.py:450: User
warnings.warn(
```



## ▼ PART(B): Logistic Regression

```
# Do as Directed:
# 5. Perform logistic regression on the admission dataset
# a) Import nyu_admission_fake.csv dataset, display it, visualize various columns

df = pd.read_csv('nyu_admission_fake.csv')
df
```

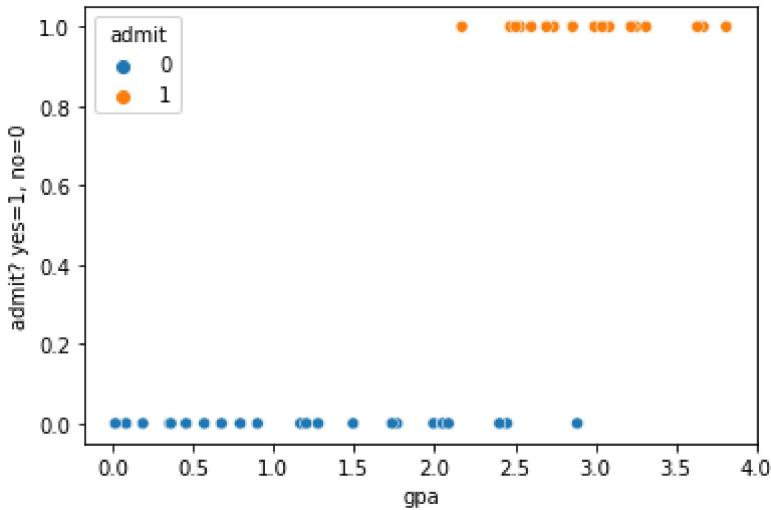


|           | <b>student</b> | <b>gpa</b> | <b>admit</b> |
|-----------|----------------|------------|--------------|
| <b>0</b>  | 0              | 3.085283   | 1            |
| <b>1</b>  | 1              | 0.083008   | 0            |
| <b>2</b>  | 2              | 2.534593   | 1            |
| <b>3</b>  | 3              | 2.995216   | 1            |
| <b>4</b>  | 4              | 1.994028   | 0            |
| <b>5</b>  | 5              | 0.899187   | 0            |
| <b>6</b>  | 6              | 0.792251   | 0            |
| <b>7</b>  | 7              | 3.042123   | 1            |
| <b>8</b>  | 8              | 0.676443   | 0            |
| <b>9</b>  | 9              | 0.353359   | 0            |
| <b>10</b> | 10             | 2.741439   | 1            |
| <b>11</b> | 11             | 3.813573   | 1            |
| <b>12</b> | 12             | 0.015793   | 0            |
| <b>13</b> | 13             | 2.048769   | 0            |
| <b>14</b> | 14             | 3.250484   | 1            |
| <b>15</b> | 15             | 2.450104   | 0            |
| <b>16</b> | 16             | 2.887021   | 0            |
| <b>17</b> | 17             | 1.167504   | 0            |
| <b>18</b> | 18             | 3.671096   | 1            |
| <b>19</b> | 19             | 2.858303   | 1            |
| <b>20</b> | 20             | 2.170177   | 1            |
| <b>21</b> | 21             | 0.568680   | 0            |
| <b>22</b> | 22             | 1.493363   | 0            |
| <b>23</b> | 23             | 2.696534   | 1            |
| <b>24</b> | 24             | 1.767333   | 0            |
| <b>25</b> | 25             | 1.736056   | 0            |
| <b>26</b> | 26             | 2.471068   | 1            |
| <b>27</b> | 27             | 2.052553   | 0            |
| <b>28</b> | 28             | 2.601589   | 1            |
| <b>29</b> | 29             | 2.404156   | 0            |
| <b>30</b> | 30             | 2.222222   | 1            |

# b) Plot the dataset on gpa vs. admit score.

```
import statsmodels.formula.api as smf

sns.scatterplot(x='gpa', y='admit', hue='admit', data=df)
plt.xlabel("gpa")
plt.ylabel("admit? yes=1, no=0")
plt.show()
fig=plt.gcf()
```



<Figure size 432x288 with 0 Axes>

# c) Find the slope and intercept of the line to fit.

```
lr=smf.ols(formula="admit ~ gpa", data=df).fit()
sns.regplot(x='gpa',y='admit', scatter=False,data=df, color='gray')
sns.scatterplot(x='gpa', y='admit', hue='admit', data=df)
plt.xlabel("gpa")
plt.ylabel("admit? yes=1, no=0")
plt.show()
fig=plt.gcf()
```

```
print(lr.params)
```

1.25 | admit |

```
# d) Compute the log odds for each entry.  
#     Merge the results with the data as a new column.  
  
# df.drop('log_odds', inplace=True, axis=1)  
# df.drop('P', inplace=True, axis=1)  
  
df['log_odds']=2.0*df['gpa']-3  
df
```



|   | student | gpa      | admit | log_odds  |
|---|---------|----------|-------|-----------|
| 0 | 0       | 3.085283 | 1     | 3.170566  |
| 1 | 1       | 0.083008 | 0     | -2.833984 |
| 2 | 2       | 2.534593 | 1     | 2.069186  |
| 3 | 3       | 2.995216 | 1     | 2.990432  |

```
# e) Using the log odds compute the probability for each entry.
```

```
df['pred_p']=np.exp(df['log_odds'])/(1+np.exp(df['log_odds']))  
df
```

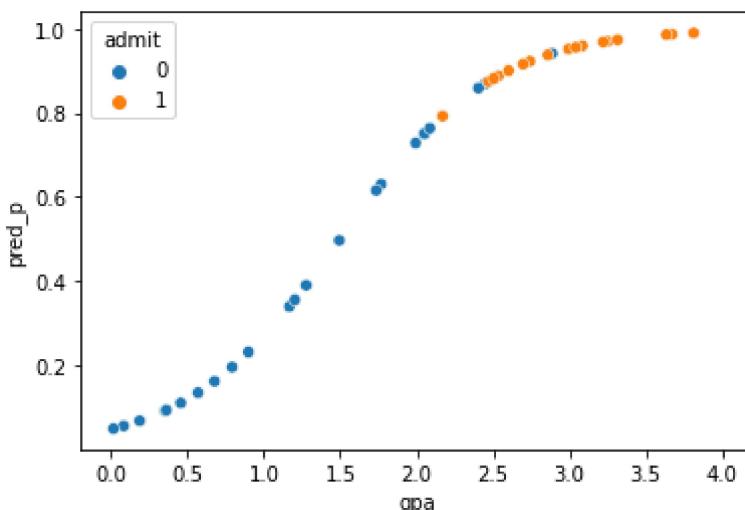


|   | student | gpa      | admit | log_odds  | pred_p   |
|---|---------|----------|-------|-----------|----------|
| 0 | 0       | 3.085283 | 1     | 3.170566  | 0.959711 |
| 1 | 1       | 0.083008 | 0     | -2.833984 | 0.055515 |
| 2 | 2       | 2.534593 | 1     | 2.069186  | 0.887872 |
| 3 | 3       | 2.995216 | 1     | 2.990432  | 0.952140 |
| 4 | 4       | 1.994028 | 0     | 0.988056  | 0.728704 |
| 5 | 5       | 0.899187 | 0     | -1.201626 | 0.231186 |
| - | -       | -        | -     | -         | -        |

# f) Plot the probabilities vs gpa graph.

```
x = np.linspace(0.0,4.0)
y = 1/(1+np.exp(-2.0*x+3))
sns.scatterplot(x='gpa',y='pred_p',hue='admit',data=df)
plt.plot(x,y,'k--',linewidth=0)
```

[<matplotlib.lines.Line2D at 0x216101f1130>]



20 20 2.170177 1 1.340354 0.792548

# g) Show the decision boundary of the regression model.

```
from sklearn.linear_model import LogisticRegression
from sklearn.datasets import load_iris

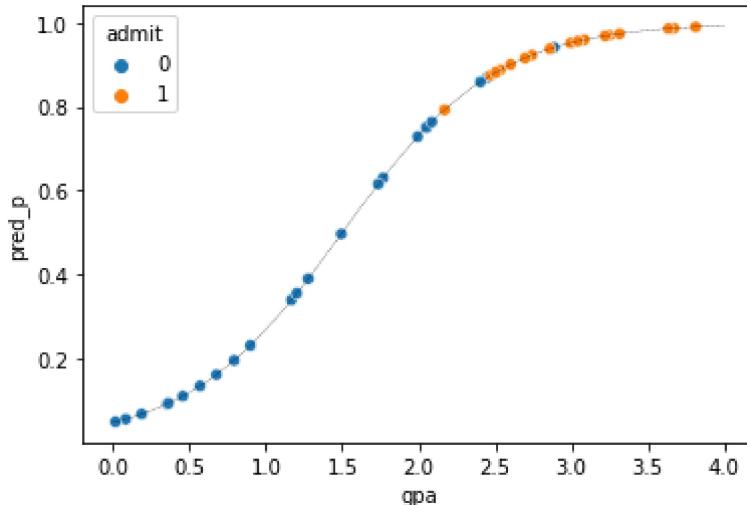
X = np.array(df['gpa'].tolist()).reshape(-1,1)
y = np.array(df['admit'].tolist())

x1_train, x1_test, y1_train, y1_test = train_test_split(X, y, train_size=0.7)

logreg = LogisticRegression().fit(x1_train,y1_train)
logreg.fit(X,y)

x = np.linspace(0.0,4.0)
y = 1/(1+np.exp(-2.0*x+3))
sns.scatterplot(x='gpa',y='pred_p',hue='admit',data=df)
plt.plot(x,y,'k--',linewidth=0.3)
```

```
[<matplotlib.lines.Line2D at 0x216115f2f70>]
```



```
# h) Show the accuracy of the regressor model.
```

```
print("Training set score: {:.3f}".format(logreg.score(x1_train,y1_train)))
print("Test set score: {:.3f}".format(logreg.score(x1_test,y1_test)))
```

```
Training set score: 0.893
Test set score: 0.917
```

```
# 6. Perform logistic regression on the credit card dataset
```

```
# a) Download credit card fraud detection dataset using
# https://www.kaggle.com/datasets/mlg-ulb/creditcardfraud?resource=downloadb)
```

```
# b) Load the dataset, visualize it, show the data headers
```

```
df=pd.read_csv('creditcard.csv')
print("shape = ",df.shape)
df
```

```
shape = (284807, 31)
```

|   | Time | V1        | V2        | V3       | V4       | V5        | V6        |       |
|---|------|-----------|-----------|----------|----------|-----------|-----------|-------|
| 0 | 0.0  | -1.359807 | -0.072781 | 2.536347 | 1.378155 | -0.338321 | 0.462388  | 0.23  |
| 1 | 0.0  | 1.191857  | 0.266151  | 0.166480 | 0.448154 | 0.060018  | -0.082361 | -0.07 |

```
display("The column headers: ")
```

```
display(df.columns.values)
```

```
'The column headers: '
```

```
array(['Time', 'V1', 'V2', 'V3', 'V4', 'V5', 'V6', 'V7', 'V8', 'V9',
       'V10', 'V11', 'V12', 'V13', 'V14', 'V15', 'V16', 'V17', 'V18',
       'V19', 'V20', 'V21', 'V22', 'V23', 'V24', 'V25', 'V26', 'V27',
       'V28', 'Amount', 'Class'], dtype=object)
```

```
# c) Preprocess the dataset if required
```

```
# i. Check duplicate data. If found remove it.
```

```
df2 = df.drop_duplicates()
```

```
df2.shape
```

```
(283726, 31)
```

```
# ii. Remove such columns which are not important. For example: Time
```

```
df2.drop(['Time'], axis=1)
```

|        | V1         | V2        | V3        | V4        | V5        | V6        | V7        |       |
|--------|------------|-----------|-----------|-----------|-----------|-----------|-----------|-------|
| 0      | -1.359807  | -0.072781 | 2.536347  | 1.378155  | -0.338321 | 0.462388  | 0.239599  | 0.09  |
| 1      | 1.191857   | 0.266151  | 0.166480  | 0.448154  | 0.060018  | -0.082361 | -0.078803 | 0.08  |
| 2      | -1.358354  | -1.340163 | 1.773209  | 0.379780  | -0.503198 | 1.800499  | 0.791461  | 0.24  |
| 3      | -0.966272  | -0.185226 | 1.792993  | -0.863291 | -0.010309 | 1.247203  | 0.237609  | 0.31  |
| 4      | -1.158233  | 0.877737  | 1.548718  | 0.403034  | -0.407193 | 0.095921  | 0.592941  | -0.21 |
| ...    | ...        | ...       | ...       | ...       | ...       | ...       | ...       | ...   |
| 284802 | -11.881118 | 10.071785 | -9.834783 | -2.066656 | -5.364473 | -2.606837 | -4.918215 | 7.30  |
| 284803 | -0.732789  | -0.055080 | 2.035030  | -0.738589 | 0.868229  | 1.058415  | 0.024330  | 0.29  |
| 284804 | 1.919565   | -0.301254 | -3.249640 | -0.557828 | 2.630515  | 3.031260  | -0.296827 | 0.70  |
| 284805 | -0.240440  | 0.530483  | 0.702510  | 0.689799  | -0.377961 | 0.623708  | -0.686180 | 0.61  |
| 284806 | -0.533413  | -0.189733 | 0.703337  | -0.506271 | -0.012546 | -0.649617 | 1.577006  | -0.41 |

```
283726 rows × 30 columns
```

```
# iii. Separate the dataset into feature column and target column. The class
# column is the target column and everything else is a feature
```

```
X=df2[['Time', 'V1', 'V2', 'V3', 'V4', 'V5', 'V6', 'V7', 'V8', 'V9',
       'V10', 'V11', 'V12', 'V13', 'V14', 'V15', 'V16', 'V17', 'V18',
       'V19', 'V20', 'V21', 'V22', 'V23', 'V24', 'V25', 'V26', 'V27',
       'V28', 'Amount']]
y=df2['Class']

display(X)
display(y)
```

|        | Time     | V1         | V2        | V3        | V4        | V5        | V6        |       |
|--------|----------|------------|-----------|-----------|-----------|-----------|-----------|-------|
| 0      | 0.0      | -1.359807  | -0.072781 | 2.536347  | 1.378155  | -0.338321 | 0.462388  | 0.23  |
| 1      | 0.0      | 1.191857   | 0.266151  | 0.166480  | 0.448154  | 0.060018  | -0.082361 | -0.07 |
| 2      | 1.0      | -1.358354  | -1.340163 | 1.773209  | 0.379780  | -0.503198 | 1.800499  | 0.79  |
| 3      | 1.0      | -0.966272  | -0.185226 | 1.792993  | -0.863291 | -0.010309 | 1.247203  | 0.23  |
| 4      | 2.0      | -1.158233  | 0.877737  | 1.548718  | 0.403034  | -0.407193 | 0.095921  | 0.59  |
| ...    | ...      | ...        | ...       | ...       | ...       | ...       | ...       | ...   |
| 284802 | 172786.0 | -11.881118 | 10.071785 | -9.834783 | -2.066656 | -5.364473 | -2.606837 | -4.91 |
| 284803 | 172787.0 | -0.732789  | -0.055080 | 2.035030  | -0.738589 | 0.868229  | 1.058415  | 0.02  |
| 284804 | 172788.0 | 1.919565   | -0.301254 | -3.249640 | -0.557828 | 2.630515  | 3.031260  | -0.29 |
| 284805 | 172788.0 | -0.240440  | 0.530483  | 0.702510  | 0.689799  | -0.377961 | 0.623708  | -0.68 |
| 284806 | 172792.0 | -0.533413  | -0.189733 | 0.703337  | -0.506271 | -0.012546 | -0.649617 | 1.57  |

283726 rows × 30 columns

```
0      0
1      0
2      0
3      0
4      0
 ..
284802  0
284803  0
284804  0
284805  0
284806  0
Name: Class, Length: 283726, dtype: int64
```

# iv. Scale the dataset using standard scaling mechanism.

```
from numpy import asarray
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
# transform data
df2 = scaler.fit_transform(df2)
print(df2)

[-1.99682292 -0.70108232 -0.04168726 ... -0.06584955  0.24419951]
```

```
-0.04086423]
[-1.9962292  0.60879165  0.16413764 ...  0.0432187 -0.34258399
-0.04086423]
[-1.99680186 -0.7003364 -0.81133678 ... -0.18382429  1.15889967
-0.04086423]

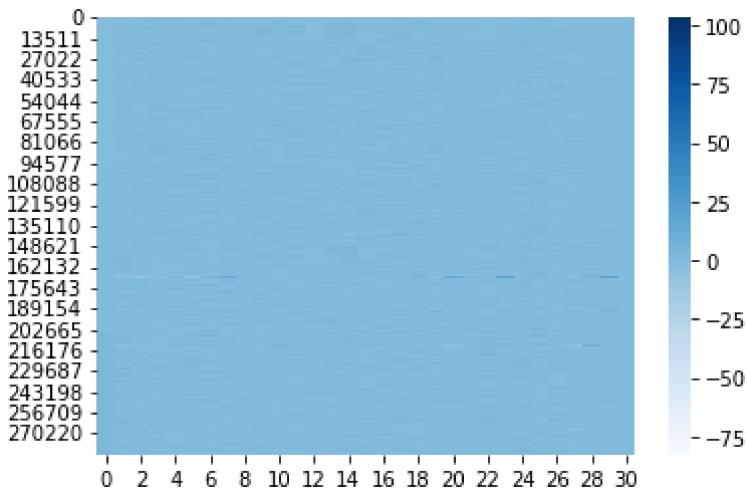
...
[ 1.64227757  0.98235398 -0.18043304 ... -0.08264021 -0.0822395
-0.04086423]
[ 1.64227757 -0.12646526  0.32465977 ...  0.31700384 -0.31339058
-0.04086423]
[ 1.64236181 -0.27686005 -0.1127094 ...  0.03994074  0.51329005
-0.04086423]]
```

```
# v. Partition dataset into training and testing set as 80%-20%
```

```
from sklearn.model_selection import train_test_split
np.random.seed(123)
X_train, X_test, y_train, y_test = train_test_split(X, y, train_size = 0.80, test_size = 0
```

```
# d) Plot histograms/heatmaps to understand the values of each variable.
```

```
sns.heatmap(df2,cmap='Blues')
plt.show()
```



```
# e) Train the model using logistic regression.
```

```
from sklearn.linear_model import LogisticRegression
lr = LogisticRegression()
lr.fit(X_train, y_train)
```

```
C:\Users\DELL\AppData\Roaming\Python\Python39\site-packages\sklearn\linear_model\log  
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.
```

# f) Obtain the training accuracy.

```
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report, roc_a  
import scikitplot as skplt
```

```
pred = lr.predict(X_train)  
matrix = confusion_matrix(y_train, pred)  
print("Training Accuracy is : ")  
accuracy_score(y_train, pred)
```

Training Accuracy is : 0.9989117983963345

# g) Test the model. Obtain the testing accuracy.

```
pred = lr.predict(X_test)  
matrix = confusion_matrix(y_test, pred)  
print("Testing Accuracy is : ")  
accuracy_score(y_test, pred)
```

Testing Accuracy is :  
0.9990836358509851

# h) Generate confusion matrix, precision and recall based on TP, FP, FN, TN.

```
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report, roc_a  
from sklearn.linear_model import LogisticRegression  
lr = LogisticRegression()  
lr.fit(X_train, y_train)
```

```
def PrintStats(cmat, y_test, pred):  
    tpos = cmat[0][0]  
    fneg = cmat[1][1]  
    fpos = cmat[0][1]  
    tneg = cmat[1][0]
```

```
lr.fit(X_train, y_train.values.ravel())  
pred = lr.predict(X_test)  
cmat = confusion_matrix(y_test, pred)  
skplt.metrics.plot_confusion_matrix(y_test, pred)  
print(classification_report(y_test, pred))  
PrintStats(cmat, y_test, pred)
```

```
C:\Users\DELL\AppData\Roaming\Python\Python39\site-packages\sklearn\linear_model\_log
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

n\_iter\_i = \_check\_optimize\_result()

```
C:\Users\DELL\AppData\Roaming\Python\Python39\site-packages\sklearn\linear_model\_log
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

n\_iter\_i = \_check\_optimize\_result()

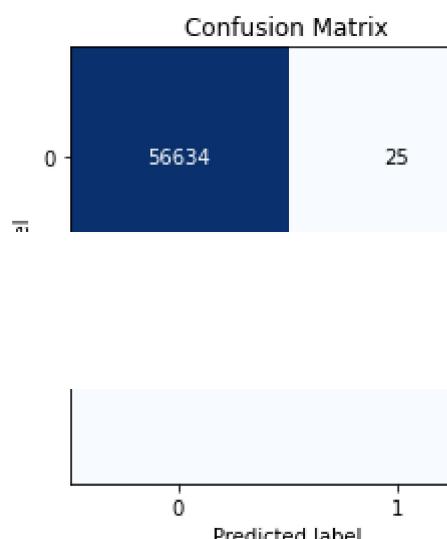
|  | precision | recall | f1-score | support |
|--|-----------|--------|----------|---------|
|--|-----------|--------|----------|---------|

|   |      |      |      |       |
|---|------|------|------|-------|
| 0 | 1.00 | 1.00 | 1.00 | 56659 |
| 1 | 0.71 | 0.69 | 0.70 | 87    |

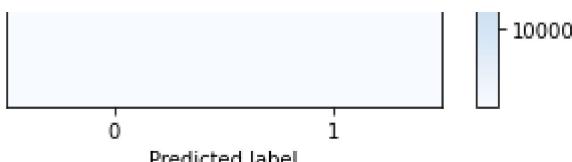
|          |  |  |      |       |
|----------|--|--|------|-------|
| accuracy |  |  | 1.00 | 56746 |
|----------|--|--|------|-------|

|           |      |      |      |       |
|-----------|------|------|------|-------|
| macro avg | 0.85 | 0.84 | 0.85 | 56746 |
|-----------|------|------|------|-------|

|              |      |      |      |       |
|--------------|------|------|------|-------|
| weighted avg | 1.00 | 1.00 | 1.00 | 56746 |
|--------------|------|------|------|-------|



-----XXXX-----



[Colab paid products](#) - [Cancel contracts here](#)

---

